

# Embedded Systems Programming

## Lecture 7


Verónica Gaspes  
`www2.hh.se/staff/vero`



CENTER FOR RESEARCH ON EMBEDDED SYSTEMS  
School of Information Science, Computer and Electrical Engineering

# Real Time?

In what ways can a program be related to time in the environment (the *real time*)?



# Real Time could be ...

An external process to sample

Reading a real-time clock is like sampling any other external process value!

An external process to react to

A program can let certain points in time denote events (e.g. by means of interrupts of a clock) to which it has to react in some way or another.

An external process to be constrained by

A program might be required to hurry enough so that some externally visible action can be performed before a certain point in time (coming later on in the course)

# Real Time could be ...

## An external process to sample

Reading a real-time clock is like sampling any other external process value!

## An external process to react to

A program can let certain points in time denote events (e.g. by means of interrupts of a clock) to which it has to react in some way or another.

## An external process to be constrained by

A program might be required to hurry enough so that some externally visible action can be performed before a certain point in time (coming later on in the course)

# Real Time could be ...

## An external process to sample

Reading a real-time clock is like sampling any other external process value!

## An external process to react to

A program can let certain points in time denote events (e.g. by means of interrupts of a clock) to which it has to react in some way or another.

## An external process to be constrained by

A program might be required to hurry enough so that some externally visible action can be performed before a certain point in time (coming later on in the course)

# Real Time could be ...

## An external process to sample

Reading a real-time clock is like sampling any other external process value!

## An external process to react to

A program can let certain points in time denote events (e.g. by means of interrupts of a clock) to which it has to react in some way or another.

## An external process to be constrained by

A program might be required to hurry enough so that some externally visible action can be performed before a certain point in time (coming later on in the course)

# Sampling the time

Requires a hardware clock that can be read like a regular external device.

## Multitude of alternatives

- Units? Seconds? Milliseconds? Cpu cycles?
- Since when? Program start? System boot? Jan 1 1970?
- Real time? Time that stops when other threads are running?  
Time that stops when CPU sleeps? Time that cannot be set  
and always increases?

# Sampling the time

Requires a hardware clock that can be read like a regular external device.

## Multitude of alternatives

- **Units?** Seconds? Milliseconds? Cpu cycles?
- **Since when?** Program start? System boot? Jan 1 1970?
- **Real time?** Time that stops when other threads are running?  
Time that stops when CPU sleeps? Time that cannot be set and always increases?



# Sampling the time

Requires a hardware clock that can be read like a regular external device.

## Multitude of alternatives

- **Units?** Seconds? Milliseconds? Cpu cycles?
- **Since when?** Program start? System boot? Jan 1 1970?
- **Real time?** Time that stops when other threads are running?  
Time that stops when CPU sleeps? Time that cannot be set and always increases?

# Sampling the time

Requires a hardware clock that can be read like a regular external device.

## Multitude of alternatives

- **Units?** Seconds? Milliseconds? Cpu cycles?
- **Since when?** Program start? System boot? Jan 1 1970?
- **Real time?** Time that stops when other threads are running?  
Time that stops when CPU sleeps? Time that cannot be set and always increases?

# Sampling the time

Requires a hardware clock that can be read like a regular external device.

## Multitude of alternatives

- **Units?** Seconds? Milliseconds? Cpu cycles?
- **Since when?** Program start? System boot? Jan 1 1970?
- **Real time?** Time that stops when other threads are running?  
Time that stops when CPU sleeps? Time that cannot be set and always increases?

# Timer/Counter1 on the AVR

What about the 16-bit counter (accessible through register TCNT1)?

## Units

CPU clock (8Mhz) divided by a programmable prescaling value (1, 8, 64, 256, 1024).

## Since when

System reset, timer reset or timer overflow (whichever was last).

## Real time

Shows real time although it can be stopped.

Aligning TCNT1 with calendar time will require calculations and extra storage (for counting overflows).

# Timer/Counter1 on the AVR

What about the 16-bit counter (accessible through register TCNT1)?

## Units

CPU clock (8Mhz) divided by a programmable prescaling value (1, 8, 64, 256, 1024).

## Since when

System reset, timer reset or timer overflow (whichever was last).

## Real time

Shows real time although it can be stopped.

Aligning TCNT1 with calendar time will require calculations and extra storage (for counting overflows).

# Timer/Counter1 on the AVR

What about the 16-bit counter (accessible through register TCNT1)?

## Units

CPU clock (8Mhz) divided by a programmable prescaling value (1, 8, 64, 256, 1024).

## Since when

System reset, timer reset or timer overflow (whichever was last).

## Real time

Shows real time although it can be stopped.

Aligning TCNT1 with calendar time will require calculations and extra storage (for counting overflows).

# Timer/Counter1 on the AVR

What about the 16-bit counter (accessible through register TCNT1)?

## Units

CPU clock (8Mhz) divided by a programmable prescaling value (1, 8, 64, 256, 1024).

## Since when

System reset, timer reset or timer overflow (whichever was last).

## Real time

Shows real time although it can be stopped.

Aligning TCNT1 with calendar time will require calculations and extra storage (for counting overflows).

# Timer/Counter1 on the AVR

What about the 16-bit counter (accessible through register TCNT1)?

## Units

CPU clock (8Mhz) divided by a programmable prescaling value (1, 8, 64, 256, 1024).

## Since when

System reset, timer reset or timer overflow (whichever was last).

## Real time

Shows real time although it can be stopped.

Aligning TCNT1 with calendar time will require calculations and extra storage (for counting overflows).



# Timestamps

In general clock readings become meaningful only when they are associated with other events: we are interested to sample to know when an event has occurred.

## Example

The clock showed 11:25 when the teacher left.

In program terms, associating a clock reading with some other event means doing the reading in close proximity to the event detection.



# Timestamps

In general clock readings become meaningful only when they are associated with other events: we are interested to sample to know when an event has occurred.

## Example

The clock showed 11:25 when the teacher left.

In program terms, associating a clock reading with some other event means doing the reading in close proximity to the event detection.



# Timestamps

In general clock readings become meaningful only when they are associated with other events: we are interested to sample to know when an event has occurred.

## Example

The clock showed 11:25 when the teacher left.

In program terms, associating a clock reading with some other event means doing the reading in close proximity to the event detection.



# Time spans

The difference between two timestamps is a value that is independent of the nominal clock values — it is a time span (characterized only by the clock resolution).

What could each timestamp mean?

- The time of some arbitrary program instruction?
- The beginning or end of a function call?
- The time of sending or receiving an asynchronous message?

All these are internal to a program under execution and have no meaning to an external observer!

# Time spans

The difference between two timestamps is a value that is independent of the nominal clock values — it is a time span (characterized only by the clock resolution).

## What could each timestamp mean?

- The time of some arbitrary program instruction?
- The beginning or end of a function call?
- The time of sending or receiving an asynchronous message?

All these are internal to a program under execution and have no meaning to an external observer!

# Time spans

The difference between two timestamps is a value that is independent of the nominal clock values — it is a time span (characterized only by the clock resolution).

## What could each timestamp mean?

- The time of some arbitrary program instruction?
- The beginning or end of a function call?
- The time of sending or receiving an asynchronous message?

All these are internal to a program under execution and have no meaning to an external observer!

# Time spans

The difference between two timestamps is a value that is independent of the nominal clock values — it is a time span (characterized only by the clock resolution).

## What could each timestamp mean?

- The time of some arbitrary program instruction?
- The beginning or end of a function call?
- The time of sending or receiving an asynchronous message?

All these are internal to a program under execution and have no meaning to an external observer!

# Time spans

The difference between two timestamps is a value that is independent of the nominal clock values — it is a time span (characterized only by the clock resolution).

## What could each timestamp mean?

- The time of some arbitrary program instruction?
- The beginning or end of a function call?
- The time of sending or receiving an asynchronous message?

All these are internal to a program under execution and have no meaning to an external observer!



# Time spans

The difference between two timestamps is a value that is independent of the nominal clock values — it is a time span (characterized only by the clock resolution).

## What could each timestamp mean?

- The time of some arbitrary program instruction?
- The beginning or end of a function call?
- The time of sending or receiving an asynchronous message?


All these are internal to a program under execution and have no meaning to an external observer!

# In a scheduled system

What looks like ...



might very well be ...




Close proximity is not the same as subsequent statements!

# In a scheduled system

What looks like ...



might very well be ...



Close proximity **is not the same as** subsequent statements!

# Time-stamping events

What we really want is to associate timestamps with the externally observable events that *drive* a system!

## Idea!

Read the clock **inside the interrupt handler** that detects the associated event

- Other interrupts are disabled while the CPU runs an interrupt handler, hence no scheduling of threads might interfere!
- There is a tight upper bound on the timestamp error, which can be calculated from CPU data and speed!

## Example

One could use an arg of an interrupt driven method for passing a timestamp!

# Time-stamping events

What we really want is to associate timestamps with the externally observable events that *drive* a system!

## Idea!

Read the clock **inside the interrupt handler** that detects the associated event

- Other interrupts are disabled while the CPU runs an interrupt handler, hence no scheduling of threads might interfere!
- There is a tight upper bound on the timestamp error, which can be calculated from CPU data and speed!

## Example

One could use an arg of an interrupt driven method for passing a timestamp!

# Time-stamping events

What we really want is to associate timestamps with the externally observable events that *drive* a system!

## Idea!

Read the clock **inside the interrupt handler** that detects the associated event

- Other interrupts are disabled while the CPU runs an interrupt handler, hence no scheduling of threads might interfere!
- There is a tight upper bound on the timestamp error, which can be calculated from CPU data and speed!

## Example

One could use an arg of an interrupt driven method for passing a timestamp!

# Time-stamping events

What we really want is to associate timestamps with the externally observable events that *drive* a system!

## Idea!

Read the clock **inside the interrupt handler** that detects the associated event

- Other interrupts are disabled while the CPU runs an interrupt handler, hence no scheduling of threads might interfere!
- There is a tight upper bound on the timestamp error, which can be calculated from CPU data and speed!

## Example

One could use an arg of an interrupt driven method for passing a timestamp!

# Time-stamping events

What we really want is to associate timestamps with the externally observable events that *drive* a system!

## Idea!

Read the clock **inside the interrupt handler** that detects the associated event

- Other interrupts are disabled while the CPU runs an interrupt handler, hence no scheduling of threads might interfere!
- There is a tight upper bound on the timestamp error, which can be calculated from CPU data and speed!

## Example


One could use an arg of an interrupt driven method for passing a timestamp!



# Example

## Calculate the speed

For a rotating wheel, measuring the time between two subsequent detections of a passing tap.




```
typedef struct{  
    Object super;  
    int previous;  
} Speedo;  
...  
Speedo speedo;  
Other client;  
INTERRUPT(SIG_XX,  
            ASYNC(&speedo,detect,TCNT1))
```

# Example

## Calculate the speed

For a rotating wheel, measuring the time between two subsequent detections of a passing tap.



```
typedef struct{  
    Object super;  
    int previous;  
} Speedo;  
...  
Speedo speedo;  
Other client;  
INTERRUPT(SIG_XX,  
           ASYNC(&speedo,detect,TCNT1))
```

# Example

## Calculate the speed

For a rotating wheel, measuring the time between two subsequent detections of a passing tap.

```
int detect(Speedo *self, int timestamp){
    SYNC(&client,
        newSpeed,
        PERIMETER/DIFF(timestamp,self->previous));
    self->previous=timestamp;
}
```

DIFF will have to take care of timer overflows!

# Example

## Calculate the speed

For a rotating wheel, measuring the time between two subsequent detections of a passing tap.

```
int detect(Speedo *self, int timestamp){  
    SYNC(&client,  
        newSpeed,  
        PERIMETER/DIFF(timestamp,self->previous));  
    self->previous=timestamp;  
}
```

DIFF will have to take care of timer overflows!


# Real-time events to react to

We know how to sample the real-time clock to obtain externally meaningful information about the passage of time.

Now suppose we want to take some action when a certain amount of time has passed.

## Example

The wheel is an engine crankshaft and we have to emit ignition signals to each cylinder



We would need a way to postpone program execution until certain points in the future.


# Real-time events to react to

We know how to sample the real-time clock to obtain externally meaningful information about the passage of time.

Now suppose we want to take some action when a certain amount of time has passed.

## Example

The wheel is an engine crankshaft and we have to emit ignition signals to each cylinder



We would need a way to postpone program execution until certain points in the future.

# Reacting to real time events

## Very poor man's solution

Consume a fixed amount of CPU cycles in a (silly) loop

```
int i;  
for(i=0;i<N;i++); // wait  
do_future_action();
```

## Problems

- Determine N by testing!
- N will be highly platform dependent!
- A lot of CPU cycles will simply be wasted!

# Reacting to real time events

## Very poor man's solution

Consume a fixed amount of CPU cycles in a (silly) loop

```
int i;  
for(i=0;i<N;i++); // wait  
do_future_action();
```

## Problems

- 1 Determine N by testing!
- 2 N will be highly platform dependent!
- 3 A lot of CPU cycles will simply be wasted!



# Reacting to real time events

## Very poor man's solution

Consume a fixed amount of CPU cycles in a (silly) loop

```
int i;  
for(i=0;i<N;i++); // wait  
do_future_action();
```

## Problems

- 1 Determine N by testing!
- 2 N will be highly platform dependent!
- 3 A lot of CPU cycles will simply be wasted!

# Reacting to real time events

## Very poor man's solution

Consume a fixed amount of CPU cycles in a (silly) loop

```
int i;  
for(i=0;i<N;i++); // wait  
do_future_action();
```

## Problems

- 1 Determine N by testing!
- 2 N will be highly platform dependent!
- 3 A lot of CPU cycles will simply be wasted!

# Reacting to real time events

## Very poor man's solution

Consume a fixed amount of CPU cycles in a (silly) loop

```
int i;  
for(i=0;i<N;i++); // wait  
do_future_action();
```

## Problems

- 1 Determine N by testing!
- 2 N will be highly platform dependent!
- 3 A lot of CPU cycles will simply be wasted!

# Reacting to real time events

## The nearly as poor man's solution

Configure a timer/counter with a known clock speed, and busy-wait for a suitable time increment

```
unsigned int i = TCNT1+N;  
while(TCNT1<i); // wait  
do_future_action();
```

## Problems

- ❶ Determine N by calculation so platform dependency disappears.
- ❷ Still a lot of wasted CPU!

# Reacting to real time events

## The nearly as poor man's solution

Configure a timer/counter with a known clock speed, and busy-wait for a suitable time increment

```
unsigned int i = TCNT1+N;  
while(TCNT1<i); // wait  
do_future_action();
```

## Problems

- 1 Determine N by calculation so platform dependency disappears.
- 2 Still a lot of wasted CPU!

# Reacting to real time events

## The nearly as poor man's solution

Configure a timer/counter with a known clock speed, and busy-wait for a suitable time increment

```
unsigned int i = TCNT1+N;  
while(TCNT1<i); // wait  
do_future_action();
```

## Problems

- 1 Determine N by calculation so platform dependency disappears.
- 2 Still a lot of wasted CPU!

# Reacting to real time events

## The nearly as poor man's solution

Configure a timer/counter with a known clock speed, and busy-wait for a suitable time increment

```
unsigned int i = TCNT1+N;  
while(TCNT1<i); // wait  
do_future_action();
```

## Problems

- 1 Determine N by calculation so platform dependency disappears.
- 2 Still a lot of wasted CPU!

# Reacting to real time events

## The standard solution

Use an Operating system call that *fakes* the timer increment busy-wait loop while making better use of the CPU resources

```
delay(N);    // wait (blocking OS call)
do_future_action();
```

- No platform dependency!
- No wasted CPU cycles (at the expense of complex OS internals)

Still a problem ...

... common to all solutions ...



# Reacting to real time events

## The standard solution

Use an Operating system call that *fakes* the timer increment busy-wait loop while making better use of the CPU resources

```
delay(N);    // wait (blocking OS call)
do_future_action();
```

- No platform dependency!
- No wasted CPU cycles (at the expense of complex OS internals)

Still a problem ...

... common to all solutions ...

# Reacting to real time events

## The standard solution

Use an Operating system call that *fakes* the timer increment busy-wait loop while making better use of the CPU resources

```
delay(N);    // wait (blocking OS call)
do_future_action();
```

- No platform dependency!
- No wasted CPU cycles (at the expense of complex OS internals)

Still a problem ...

... common to all solutions ...

# Reacting to real time events

## The standard solution

Use an Operating system call that *fakes* the timer increment busy-wait loop while making better use of the CPU resources

```
delay(N);    // wait (blocking OS call)
do_future_action();
```

- No platform dependency!
- No wasted CPU cycles (at the expense of complex OS internals)

Still a problem ...

... common to all solutions ...

# Reacting to real time events

## The standard solution

Use an Operating system call that *fakes* the timer increment busy-wait loop while making better use of the CPU resources

```
delay(N);    // wait (blocking OS call)
do_future_action();
```


- No platform dependency!
- No wasted CPU cycles (at the expense of complex OS internals)

## Still a problem ...

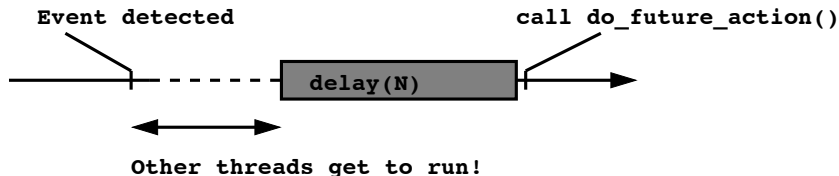
... common to all solutions ...

# In a scheduled system

What looks like ...




might very well be ...



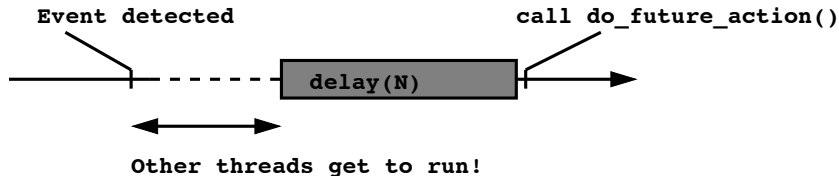
Had we known the scheduler's choice, a smaller  $N$  had been used!

# In a scheduled system

What looks like ...



might very well be ...



Had we known the scheduler's choice, a smaler N had been used!

# Relative delays

The problem is that these delay services are always specified using **relative time**:

- The constructed real-time event will occur at a time obtained by adding the delay parameter  $N$  to *now*.
- But *now* is not a very meaningful time reference in a scheduled system, as it is not related to any externally observable signals.

Other common OS services share this problem: `sleep`, `usleep` and `nanosleep`.

We are not going to use OS services in the course.

# Relative delays

The problem is that these delay services are always specified using **relative time**:

- The constructed real-time event will occur at a time obtained by adding the delay parameter  $N$  to *now*.
- But *now* is not a very meaningful time reference in a scheduled system, as it is not related to any externally observable signals.

Other common OS services share this problem: `sleep`, `usleep` and `nanosleep`.

We are not going to use OS services in the course.



# Relative delays

The problem is that these delay services are always specified using **relative time**:

- The constructed real-time event will occur at a time obtained by adding the delay parameter  $N$  to *now*.
- But *now* is not a very meaningful time reference in a scheduled system, as it is not related to any externally observable signals.

Other common OS services share this problem: `sleep`, `usleep` and `nanosleep`.

We are not going to use OS services in the course.

# Relative delays

The problem is that these delay services are always specified using **relative time**:

- The constructed real-time event will occur at a time obtained by adding the delay parameter  $N$  to *now*.
- But *now* is not a very meaningful time reference in a scheduled system, as it is not related to any externally observable signals.

Other common OS services share this problem: `sleep`, `usleep` and `nanosleep`.

We are not going to use OS services in the course.

# Relative delays

The problem is that these delay services are always specified using **relative time**:

- The constructed real-time event will occur at a time obtained by adding the delay parameter  $N$  to *now*.
- But *now* is not a very meaningful time reference in a scheduled system, as it is not related to any externally observable signals.

Other common OS services share this problem: `sleep`, `usleep` and `nanosleep`.

We are not going to use OS services in the course.

# Yet another problem

Even if other threads were not to interfere, using delay services as a means of specifying future points in time has another fundamental drawback.

## Example

Consider a task running a CPU-heavy function `do_work()` every 100 milliseconds. Using `delay()`, the naive implementation would be:

```
while(1){  
    do_work();  
    delay(100);  
}
```

## Yet another problem


Even if other threads were not to interfere, using delay services as a means of specifying future points in time has another fundamental drawback.

### Example

Consider a task running a CPU-heavy function `do_work()` every 100 milliseconds. Using `delay()`, the naive implementation would be:

```
while(1){  
    do_work();  
    delay(100);  
}
```

# Accumulating drift




X is the time take to do\_work

With relative delays, each turn in the loop will take at least  $100+X$  milliseconds.

A drift of X milliseconds will accumulate every turn!

# Accumulating drift




X is the time take to do\_work

With relative delays, each turn in the loop will take at least  $100+X$  milliseconds.

A drift of X milliseconds will accumulate every turn!

# Accumulating drift




X is the time take to do\_work

With relative delays, each turn in the loop will take at least  $100 + X$  milliseconds.

A drift of X milliseconds will accumulate every turn!




# Accumulating drift



With other threads in the system, the already bad scenario gets worse!

That means that even if we knew  $X$ , we wouldn't be able to compensate the delay time in any predictable manner!

# Accumulating drift



With other threads in the system, the already bad scenario gets worse!

That means that even if we knew  $X$ , we wouldn't be able to compensate the delay time in any predictable manner!

# A stable reference

What we need is a stable time reference to use as a basis whenever we specify a relative time (instead of now).

## Baselines


We introduce **the baseline of a message** to mean the earliest time a message is allowed to start.

## Time stamps of interrupts!

The baseline of an event is its timestamp:

**Baseline: start after**

**Actual method execution**



# A stable reference

What we need is a stable time reference to use as a basis whenever we specify a relative time (instead of now).

## Baselines


We introduce **the baseline of a message** to mean the earliest time a message is allowed to start.

Time stamps of interrupts!

The baseline of an event is its timestamp:

**Baseline: start after**

**Actual method execution**



# A stable reference

What we need is a stable time reference to use as a basis whenever we specify a relative time (instead of now).

## Baselines

We introduce **the baseline of a message** to mean the earliest time a message is allowed to start.

## Time stamps of interrupts!

The baseline of an event is its timestamp:

**Baseline: start after**

**Actual method execution**




# A stable reference

## SYNC

Calling methods with SYNC doesn't change the baseline (the call inherits the baseline)

**Baseline: start after**



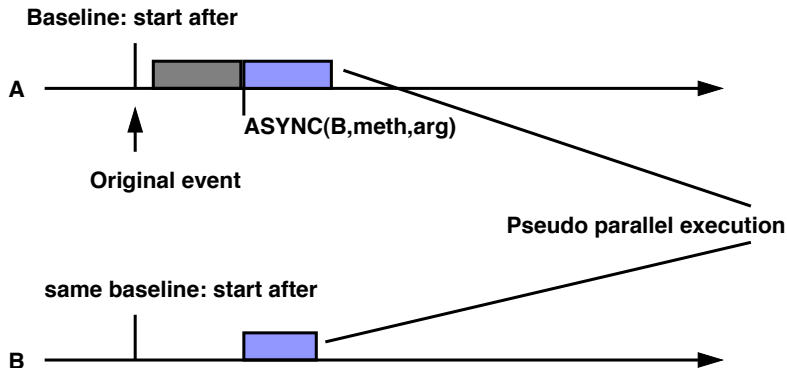
**same baseline: start after**



# A stable reference

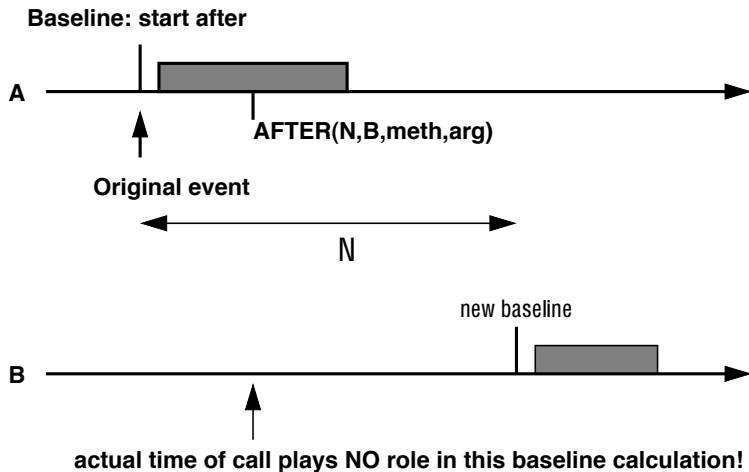
## ASYNCR

By default ASYNCR method calls will inherit the baseline



# A stable reference

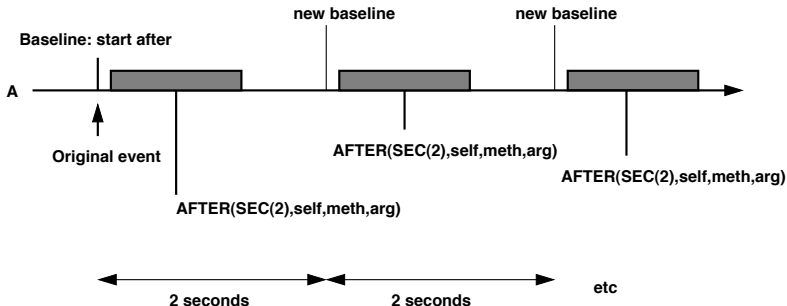
For ASYNC we may also consider adding a baseline offset  $N$ !





# Periodic tasks

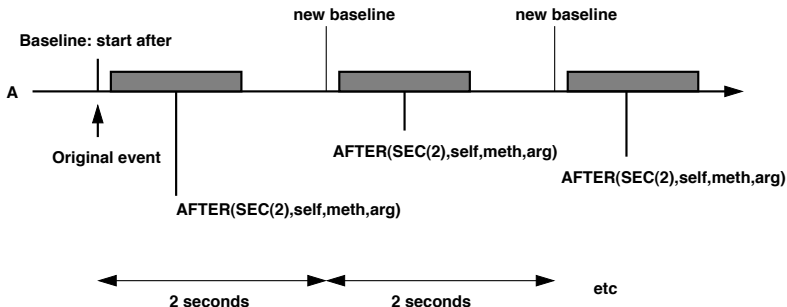
To create a cyclic reaction, simply call **self** with the same method and a new baseline:



SEC is a convenient macro that makes the call independent of current timer resolution.

# Periodic tasks

To create a cyclic reaction, simply call **self** with the same method and a new baseline:



SEC is a convenient macro that makes the call independent of current timer resolution.

# Implementing AFTER

- ① Let the baseline be stored in every message (as part of the Msg structure)
- ② AFTER is the same as ASYNC, but
  - New baseline is  
`MAX(now, offset+current->baseline)`
  - If `baseline > now`, put message in a `timerQ` instead of `readyQ`
  - Set up a timer to generate an interrupt after earliest baseline
  - At each timer interrupt, move first `timerQ` message to `readyQ` and configure a new timer interrupt

In fact ASYNC can now be defined as

```
#define ASYNC(to, meth, arg) AFTER(0, to, meth, arg)
```

# Implementing AFTER

- 1 Let the baseline be stored in every message (as part of the Msg structure)
- 2 AFTER is the same as ASYNC, but
  - New baseline is  
`MAX(now, offset+current->baseline)`
  - If `baseline > now`, put message in a `timerQ` instead of `readyQ`
  - Set up a timer to generate an interrupt after earliest baseline
  - At each timer interrupt, move first `timerQ` message to `readyQ` and configure a new timer interrupt

In fact ASYNC can now be defined as

```
#define ASYNC(to, meth, arg) AFTER(0, to, meth, arg)
```

# Implementing AFTER

- ① Let the baseline be stored in every message (as part of the Msg structure)
- ② AFTER is the same as ASYNC, but
  - New baseline is  
`MAX(now, offset+current->baseline)`
  - If `baseline > now`, put message in a **timerQ** instead of `readyQ`
  - Set up a timer to generate an interrupt after earliest baseline
  - At each timer interrupt, move first timerQ message to `readyQ` and configure a new timer interrupt

In fact ASYNC can now be defined as

```
#define ASYNC(to, meth, arg) AFTER(0, to, meth, arg)
```

# Implementing AFTER

- ① Let the baseline be stored in every message (as part of the Msg structure)
- ② AFTER is the same as ASYNC, but
  - New baseline is  
`MAX(now, offset+current->baseline)`
  - If `baseline > now`, put message in a **timerQ** instead of `readyQ`
  - Set up a timer to generate an interrupt after earliest baseline
  - At each timer interrupt, move first timerQ message to `readyQ` and configure a new timer interrupt

In fact ASYNC can now be defined as

```
#define ASYNC(to,meth,arg) AFTER(0,to,meth,arg)
```

# Implementing AFTER

- ① Let the baseline be stored in every message (as part of the Msg structure)
- ② AFTER is the same as ASYNC, but
  - New baseline is  
`MAX(now, offset+current->baseline)`
  - If `baseline > now`, put message in a **timerQ** instead of readyQ
  - Set up a timer to generate an interrupt after earliest baseline
  - At each timer interrupt, move first timerQ message to readyQ and configure a new timer interrupt

In fact ASYNC can now be defined as

```
#define ASYNC(to, meth, arg) AFTER(0, to, meth, arg)
```

# Implementing AFTER

- ① Let the baseline be stored in every message (as part of the Msg structure)
- ② AFTER is the same as ASYNC, but
  - New baseline is  
 $\text{MAX}(\text{now}, \text{offset} + \text{current} \rightarrow \text{baseline})$
  - If  $\text{baseline} > \text{now}$ , put message in a **timerQ** instead of readyQ
  - Set up a timer to generate an interrupt after earliest baseline
  - At each timer interrupt, move first timerQ message to readyQ and configure a new timer interrupt

In fact ASYNC can now be defined as

```
#define ASYNC(to, meth, arg) AFTER(0, to, meth, arg)
```



# Implementing AFTER

- ① Let the baseline be stored in every message (as part of the Msg structure)
- ② AFTER is the same as ASYNC, but
  - New baseline is  
 $\text{MAX}(\text{now}, \text{offset} + \text{current} \rightarrow \text{baseline})$
  - If  $\text{baseline} > \text{now}$ , put message in a **timerQ** instead of readyQ
  - Set up a timer to generate an interrupt after earliest baseline
  - At each timer interrupt, move first timerQ message to readyQ and configure a new timer interrupt

In fact ASYNC can now be defined as

```
#define ASYNC(to, meth, arg) AFTER(0, to, meth, arg)
```

# Implementing AFTER

- ① Let the baseline be stored in every message (as part of the Msg structure)
- ② AFTER is the same as ASYNC, but
  - New baseline is  
 $\text{MAX}(\text{now}, \text{offset} + \text{current} \rightarrow \text{baseline})$
  - If  $\text{baseline} > \text{now}$ , put message in a **timerQ** instead of readyQ
  - Set up a timer to generate an interrupt after earliest baseline
  - At each timer interrupt, move first timerQ message to readyQ and configure a new timer interrupt

In fact ASYNC can now be defined as

```
#define ASYNC(to, meth, arg) AFTER(0, to, meth, arg)
```